

Exploratory Data Analysis Starter

Import packages

```
In [1]: import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# Shows plots in jupyter notebook
%matplotlib inline

# Set plot style
sns.set(color_codes=True)
```

Loading data with Pandas

We need to load `client_data.csv` and `price_data.csv` into individual dataframes so that we can work with them in Python. For this notebook and all further notebooks, it will be assumed that the CSV files will be placed in the same file location as the notebook. If they are not, please adjust the directory within the `read_csv` method accordingly.

```
In [2]: client_df = pd.read_csv('./client_data.csv')
price_df = pd.read_csv('./price_data.csv')
```

You can view the first 3 rows of a dataframe using the `head` method. Similarly, if you wanted to see the last 3, you can use `tail(3)`

```
In [3]: client_df.head(3)
```

Out[3]:

	id	channel_sales	cons_12m	cons_gas_12m	cons_last_month	date_activ	date_end	date_modif_prod	date_renewal	forecast_cons_12m	...
0	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwkcsoibcdxkica	0	54946	0	2013-06-15	2016-06-15	2015-11-01	2015-06-23	0.00	...
1	d29c2c54acc38ff3c0614d0a653813dd	MISSING	4660	0	0	2009-08-21	2016-08-30	2009-08-21	2015-08-31	189.95	...
2	764c75f661154dac3a6c254cd082ea7d	foosdfpfkusacimwkcsoibcdxkica	544	0	0	2010-04-16	2016-04-16	2010-04-16	2015-04-17	47.96	...

3 rows × 26 columns

```
In [4]: price_df.head(3)
```

Out[4]:

	id	price_date	price_off_peak_var	price_peak_var	price_mid_peak_var	price_off_peak_fix	price_peak_fix	price_mid_peak_fix
0	038af19179925da21a25619c5a24b745	2015-01-01	0.151367	0.0	0.0	44.266931	0.0	0.0
1	038af19179925da21a25619c5a24b745	2015-02-01	0.151367	0.0	0.0	44.266931	0.0	0.0
2	038af19179925da21a25619c5a24b745	2015-03-01	0.151367	0.0	0.0	44.266931	0.0	0.0

Descriptive statistics of data

Data types

It is useful to first understand the data that you're dealing with along with the data types of each column. The data types may dictate how you transform and engineer features.

To get an overview of the data types within a data frame, use the `info()` method.

```
In [5]: client_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14606 entries, 0 to 14605
Data columns (total 26 columns):
#   Column                                Non-Null Count  Dtype
---  ---
0   id                                    14606 non-null  object
1   channel_sales                        14606 non-null  object
2   cons_12m                             14606 non-null  int64
3   cons_gas_12m                         14606 non-null  int64
4   cons_last_month                      14606 non-null  int64
5   date_activ                           14606 non-null  object
6   date_end                             14606 non-null  object
7   date_modif_prod                      14606 non-null  object
8   date_renewal                         14606 non-null  object
9   forecast_cons_12m                    14606 non-null  float64
10  forecast_cons_year                   14606 non-null  int64
11  forecast_discount_energy              14606 non-null  float64
12  forecast_meter_rent_12m               14606 non-null  float64
13  forecast_price_energy_off_peak        14606 non-null  float64
14  forecast_price_energy_peak            14606 non-null  float64
15  forecast_price_pow_off_peak           14606 non-null  float64
16  has_gas                              14606 non-null  object
17  imp_cons                             14606 non-null  float64
18  margin_gross_pow_ele                  14606 non-null  float64
19  margin_net_pow_ele                    14606 non-null  float64
20  nb_prod_act                           14606 non-null  int64
21  net_margin                           14606 non-null  float64
22  num_years_antig                       14606 non-null  int64
23  origin_up                             14606 non-null  object
24  pow_max                              14606 non-null  float64
25  churn                                14606 non-null  int64
dtypes: float64(11), int64(7), object(8)
memory usage: 2.9+ MB
```

```
In [6]: price_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 193002 entries, 0 to 193001
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   id                    193002 non-null object  
1   price_date            193002 non-null object  
2   price_off_peak_var    193002 non-null float64 
3   price_peak_var        193002 non-null float64 
4   price_mid_peak_var    193002 non-null float64 
5   price_off_peak_fix    193002 non-null float64 
6   price_peak_fix        193002 non-null float64 
7   price_mid_peak_fix    193002 non-null float64 
dtypes: float64(6), object(2)
memory usage: 11.8+ MB
```

Statistics

Now let's look at some statistics about the datasets. We can do this by using the `describe()` method.

```
In [7]: client_df.describe()

Out[7]:
```

	cons_12m	cons_gas_12m	cons_last_month	forecast_cons_12m	forecast_cons_year	forecast_discount_energy	forecast_meter_rent_12m	forecast_price_energy_off_peak	forecast_price
count	1.460600e+04	1.460600e+04	14606.000000	14606.000000	14606.000000	14606.000000	14606.000000	14606.000000	
mean	1.592203e+05	2.809238e+04	16090.269752	1868.614880	1399.762906	0.966726	63.086871	0.137283	
std	5.734653e+05	1.629731e+05	64364.196422	2387.571531	3247.786255	5.108289	66.165783	0.024623	
min	0.000000e+00	0.000000e+00	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	5.674750e+03	0.000000e+00	0.000000	494.995000	0.000000	0.000000	16.180000	0.116340	
50%	1.411550e+04	0.000000e+00	792.500000	1112.875000	314.000000	0.000000	18.795000	0.143166	
75%	4.076375e+04	0.000000e+00	3383.000000	2401.790000	1745.750000	0.000000	131.030000	0.146348	
max	6.207104e+06	4.154590e+06	771203.000000	82902.830000	175375.000000	30.000000	599.310000	0.273963	

```
In [8]: price_df.describe()

Out[8]:
```

	price_off_peak_var	price_peak_var	price_mid_peak_var	price_off_peak_fix	price_peak_fix	price_mid_peak_fix
count	193002.000000	193002.000000	193002.000000	193002.000000	193002.000000	193002.000000
mean	0.141027	0.054630	0.030496	43.334477	10.622875	6.409984
std	0.025032	0.049924	0.036298	5.410297	12.841895	7.773592
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.125976	0.000000	0.000000	40.728885	0.000000	0.000000
50%	0.146033	0.085483	0.000000	44.266930	0.000000	0.000000
75%	0.151635	0.101673	0.072558	44.444710	24.339581	16.226389
max	0.280700	0.229788	0.114102	59.444710	36.490692	17.458221

Data visualization

If you're working in Python, two of the most popular packages for visualization are `matplotlib` and `seaborn`. We highly recommend you use these, or at least be familiar with them because they are ubiquitous!

Below are some functions that you can use to get started with visualizations.

```
In [9]: def plot_stacked_bars(dataframe, title_, size_=(18, 10), rot_=0, legend_="upper right"):
        """
        Plot stacked bars with annotations
        """
        ax = dataframe.plot(
            kind="bar",
            stacked=True,
            figsize=size_,
            rot=rot_,
            title=title_
        )

        # Annotate bars
        annotate_stacked_bars(ax, fontsize=14)
        # Rename Legend
        plt.legend(["Retention", "Churn"], loc=legend_)
        # Labels
        plt.ylabel("Company base (%)")
        plt.show()

def annotate_stacked_bars(ax, pad=0.99, colour="white", fontsize=13):
    """
    Add value annotations to the bars
    """

    # Iterate over the plotted rectangles/bars
    for p in ax.patches:

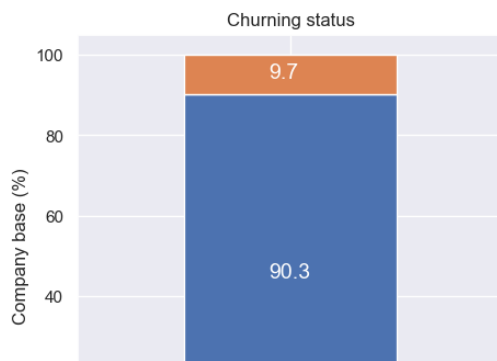
        # Calculate annotation
        value = str(round(p.get_height(),1))
        # If value is 0 do not annotate
        if value == '0.0':
            continue
        ax.annotate(
            value,
            ((p.get_x()+ p.get_width()/2)*pad-0.05, (p.get_y()+p.get_height()/2)*pad),
            color=colour,
            size=fontsize
        )

def plot_distribution(dataframe, column, ax, bins_=50):
    """
    Plot variable distribution in a stacked histogram of churned or retained company
    """

    # Create a temporal dataframe with the data to be plot
    temp = pd.DataFrame({"Retention": dataframe[dataframe["churn"]==0][column],
        "Churn":dataframe[dataframe["churn"]==1][column]})
    # Plot the histogram
    temp[["Retention", "Churn"]].plot(kind='hist', bins=bins_, ax=ax, stacked=True)
    # X-axis Label
    ax.set_xlabel(column)
    # Change the x-axis to plain style
    ax.ticklabel_format(style='plain', axis='x')
```

The first function `plot_stacked_bars` is used to plot a stacked bar chart. An example of how you could use this is shown below:

```
In [10]: churn = client_df[['id', 'churn']]
churn.columns = ['Companies', 'churn']
churn_total = churn.groupby(churn['churn']).count()
churn_percentage = churn_total / churn_total.sum() * 100
plot_stacked_bars(churn_percentage.transpose(), "Churning status", (5, 5), legend_="lower right")
```



9.6% of customers have churned

The second function `annotate_bars` is used by the first function, but the third function `plot_distribution` helps you to plot the distribution of a numeric column. An example of how it can be used is given below:




Consumption


```
In [19]: consumption = client_df[['id', 'cons_12m', 'cons_gas_12m', 'cons_last_month', 'imp_cons', 'has_gas', 'churn']]

fig, axs = plt.subplots(nrows=4, figsize=(15, 35))

plot_distribution(consumption, 'cons_12m', axs[0])
plot_distribution(consumption, 'cons_gas_12m', axs[1])
plot_distribution(consumption, 'cons_last_month', axs[2])
plot_distribution(consumption, 'imp_cons', axs[3])

plt.show()
```



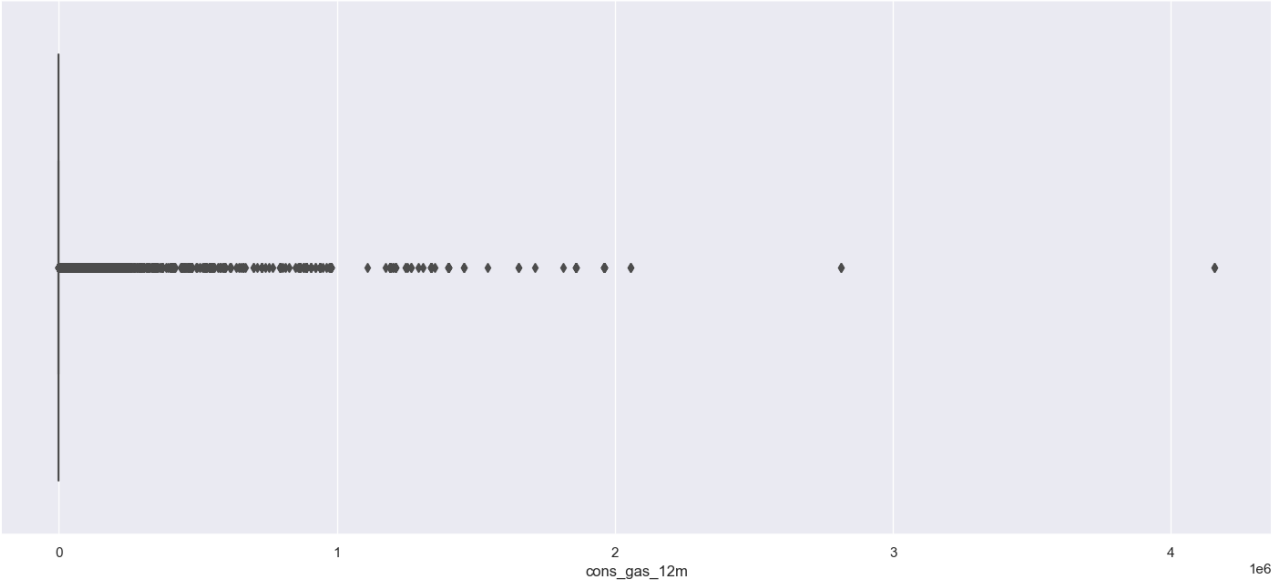




```
In [22]: consumption = client_df[['id', 'cons_12m', 'cons_gas_12m', 'cons_last_month', 'imp_cons', 'has_gas', 'churn']]

fig, axs = plt.subplots(nrows=4, figsize=(18, 35))

sns.boxplot(consumption.cons_12m, ax= axs[0])
sns.boxplot(consumption.cons_gas_12m, ax=axs[1])
sns.boxplot(consumption.cons_last_month, ax=axs[2])
sns.boxplot(consumption.imp_cons, ax=axs[3])
```

```
Out[22]: <AxesSubplot:xlabel='imp_cons'>
```


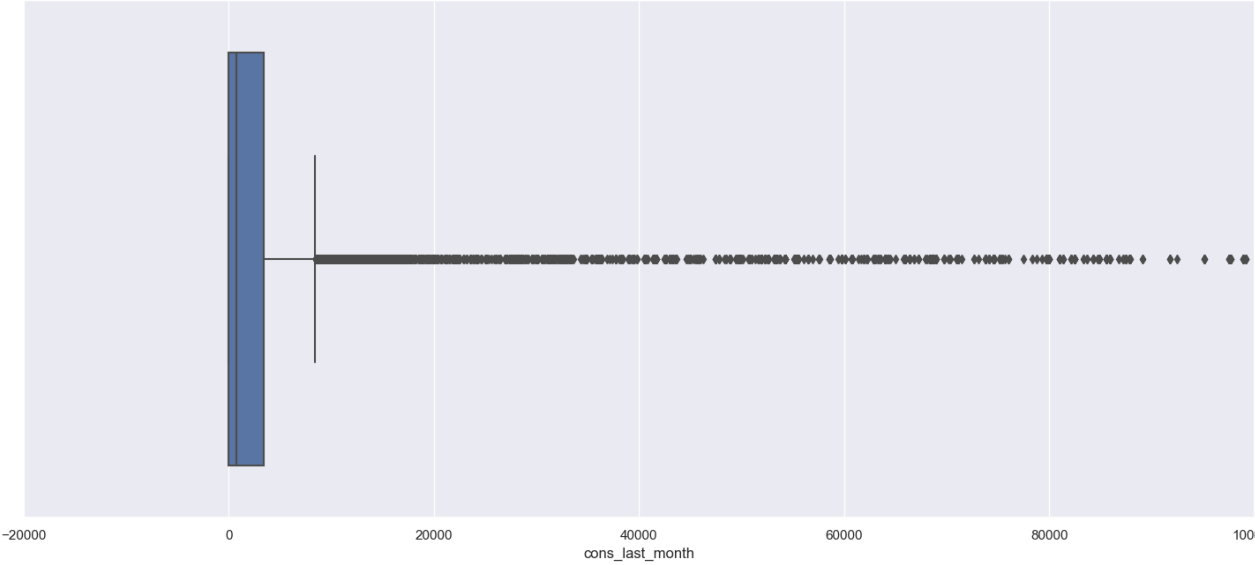
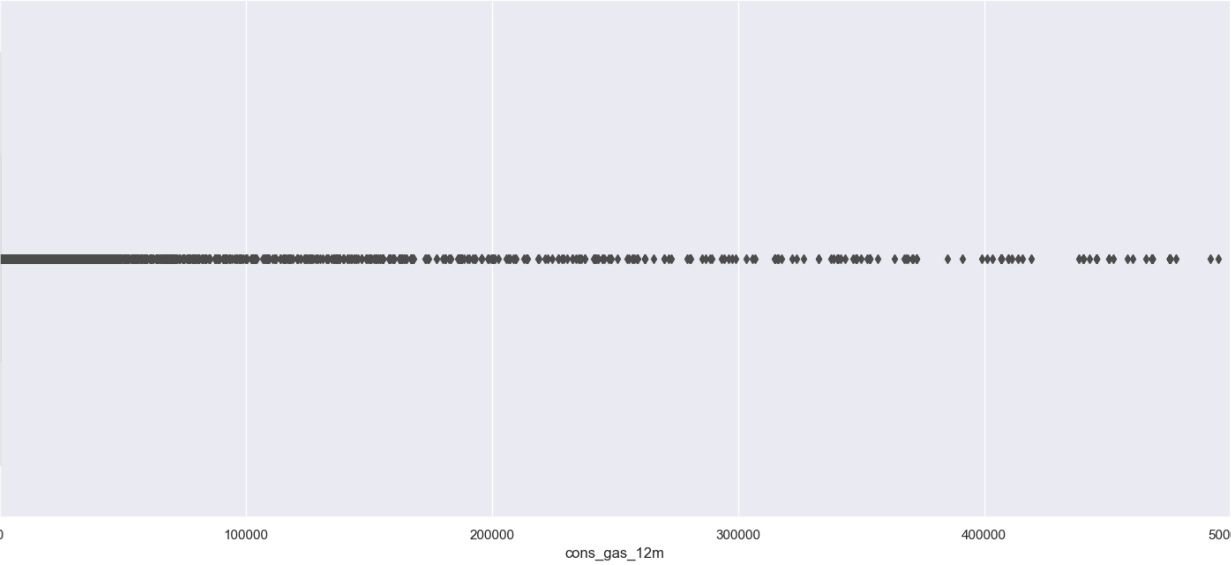




```
In [23]: consumption = client_df[['id', 'cons_12m', 'cons_gas_12m', 'cons_last_month', 'imp_cons', 'has_gas', 'churn']]

fig, axs = plt.subplots(nrows=4, figsize=(18, 35))

sns.boxplot(consumption.cons_12m, ax=axs[0])
sns.boxplot(consumption.cons_gas_12m, ax=axs[1])
sns.boxplot(consumption.cons_last_month, ax=axs[2])
sns.boxplot(consumption.imp_cons, ax=axs[3])

for ax in axs:
    ax.ticklabel_format(style='plain', axis='x')
    # Set x-axis limit
    axs[0].set_xlim(-200000, 2000000)
    axs[1].set_xlim(0, 500000)
    axs[2].set_xlim(-20000, 100000)
    plt.show()
```



Forecast

```
In [24]: forecast = client_df[
    [ "forecast_cons_12m",
      "forecast_cons_year", "forecast_discount_energy", "forecast_meter_rent_12m",
      "forecast_price_energy_off_peak", "forecast_price_energy_peak",
      "forecast_price_pow_off_peak", "churn"
    ]
]

fig, axs = plt.subplots(nrows=7, figsize=(18,50))




# Plot histogram
plot_distribution(client_df, "forecast_cons_12m", axs[0])
plot_distribution(client_df, "forecast_cons_year", axs[1])
plot_distribution(client_df, "forecast_discount_energy", axs[2])
plot_distribution(client_df, "forecast_meter_rent_12m", axs[3])
plot_distribution(client_df, "forecast_price_energy_off_peak", axs[4])
plot_distribution(client_df, "forecast_price_energy_peak", axs[5])
plot_distribution(client_df, "forecast_price_pow_off_peak", axs[6])
```






```
In [25]: margin = client_df[['margin_gross_pow_ele', 'margin_net_pow_ele', 'net_margin', 'churn']]
margin.head()
```


Out[25]:

	margin_gross_pow_ele	margin_net_pow_ele	net_margin	churn
0	25.44	25.44	678.99	1
1	16.38	16.38	18.89	0
2	28.60	28.60	6.60	0
3	30.22	30.22	25.46	0
4	44.91	44.91	47.98	0






```
In [26]: fig, axs = plt.subplots(nrows=3, figsize=(18,50))  
  
# Plot histogram  
plot_distribution(client_df, "margin_gross_pow_ele", axs[0])  
plot_distribution(client_df, "margin_net_pow_ele", axs[1])  
plot_distribution(client_df, "net_margin", axs[2])
```



```
In [27]: fig, axs = plt.subplots(nrows=3, figsize=(18, 35))
sns.boxplot(client_df["margin_gross_pow_ele"], ax=axs[0])
sns.boxplot(client_df["margin_net_pow_ele"], ax=axs[1])
sns.boxplot(client_df["net_margin"], ax=axs[2])
```

```
Out[27]: <AxesSubplot:xlabel='net_margin'>
```







plot the heatmap

```
In [33]: corr = client_df.corr()
plt.figure(figsize=(15,10))
sns.heatmap(corr, xticklabels=corr.columns, yticklabels=corr.columns, annot=True, cmap='Blues',fmt=".2f" )
```

```
Out[33]: <AxesSubplot:>
```



- cons_last_month and cons_12m has high correlation
- forecast_cons_year and forecast_cons_12m has high correlation
- imp_cons and forecast_cons_year have high correlation
- forecast_cons_12m and net_margin have high correlation

Findings:

* Analyzed consumption, forecast, margin and power related columns from client data

* We have highly positively skewed data in almost all areas. This needs to be properly handled before doing data modelling

* Analysis shows that we have around 9.7% of customers who have churned

Suggestions:

* Churning is likely to happen when a competitor has given good offers at the same price as here

* We can ask for customer feedback, to check for the suggestions/complaints from their side

* We can provide extra benefits and offers for people who subscribed with the company for a long/specified period of time. This can help to reduce the churn percentage